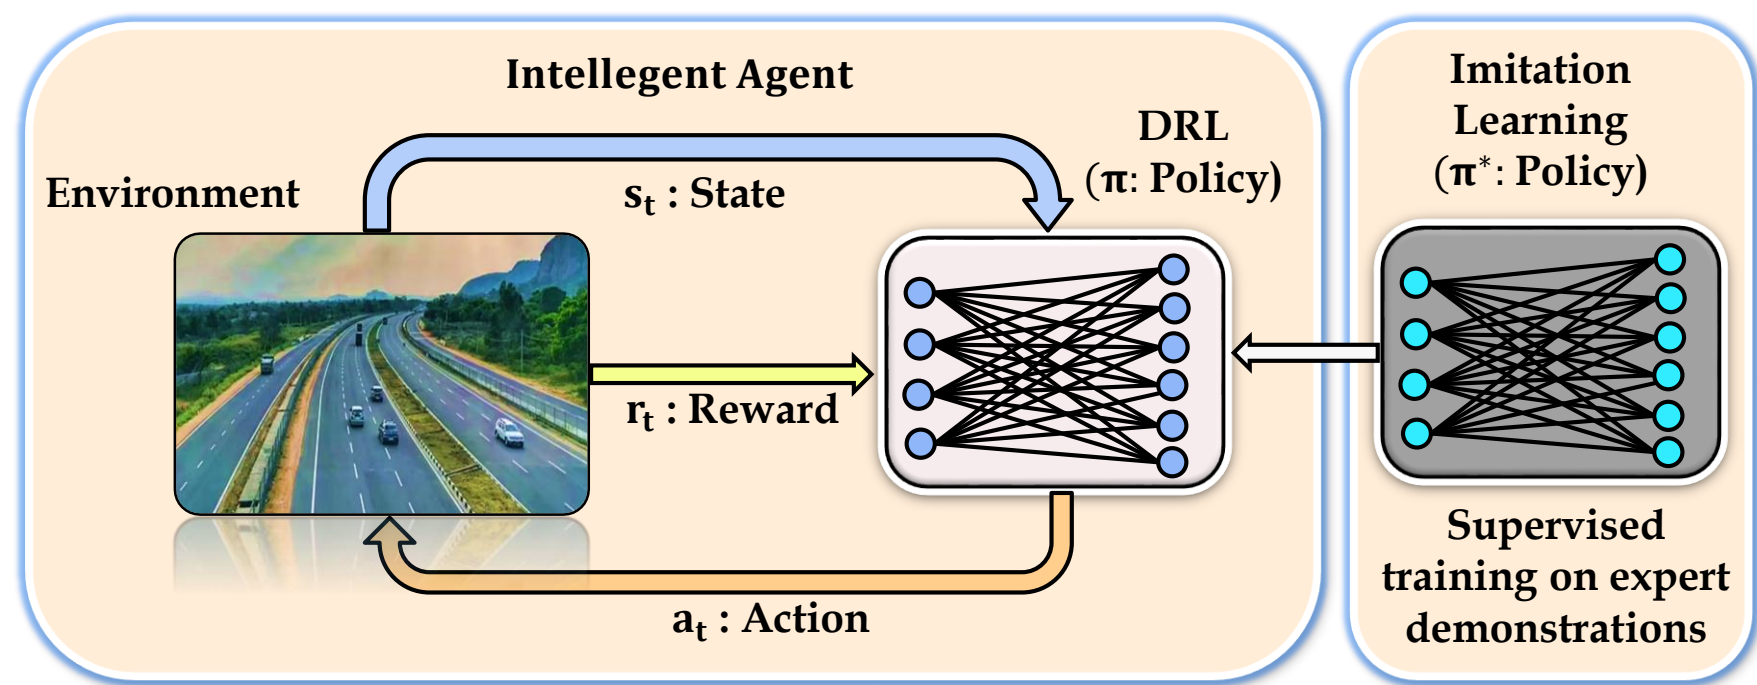


## Problem Statement

- Vision-based motion planning is a crucial task for an intelligent agent to effectively learn and adapt to the dynamic characteristics of its environment.



### Challenges in IL-Based DRL Approaches:


- Performance limited by the quality of the expert data.
- Impossible to collect expert data for all driving scenarios.
- Distribution mismatch between training and real-world data causes poor generalization to new states.

## Applications



## Motivation

- Greater Road Safety
- Greater Independence
- Reduced Congestion
- Increase Productivity
- Enhanced accessibility



- According to US market reports [5], transportation applications dominate the AV market with 65% share across self-driving cars, trucks, and public transit. This dominance creates significant opportunities for AI-driven solutions in autonomous driving.

## Key Contributions

- We propose a Diffusion-Guided Deep Reinforcement Learning (DGDRL) framework, which integrates diffusion models with Soft Actor-Critic to enhance task generalization and reduce dependency on IL.
- We introduce a novel modified Partially Observable Markov Decision Process (mPOMDP) to optimize the behavior of the SAC policy for handling uncertainty in dynamic environments.
- We conduct extensive testing across multiple CARLA towns under varying weather conditions and traffic densities (Empty, Regular, and Dense scenarios), where DGDRL achieves 95% task completion success rate with enhanced safety

### Actor Network Function ( $\pi$ )

$$L_{\phi}^{\pi} = -\mathbb{E}_{o_t \sim D} \left[ \mathbb{E}_{\tilde{a}_t \sim \pi_{\phi}(\cdot|o_t)} \left( \min_{k=1,2} Q_{\theta_k}(o_t, \tilde{a}_t) - \alpha \log \pi_{\phi}(\tilde{a}_t|o_t) \right) \right] + \lambda \mathbb{E}_{o'_t \sim \text{Diffusion}(o_t)} \left[ D_{\text{KL}}(\pi_{\phi}(\cdot|o_t) \parallel \pi_{\phi}(\cdot|o'_t)) \right]$$

## Proposed Method : Diffusion-Guided Deep Reinforcement Learning

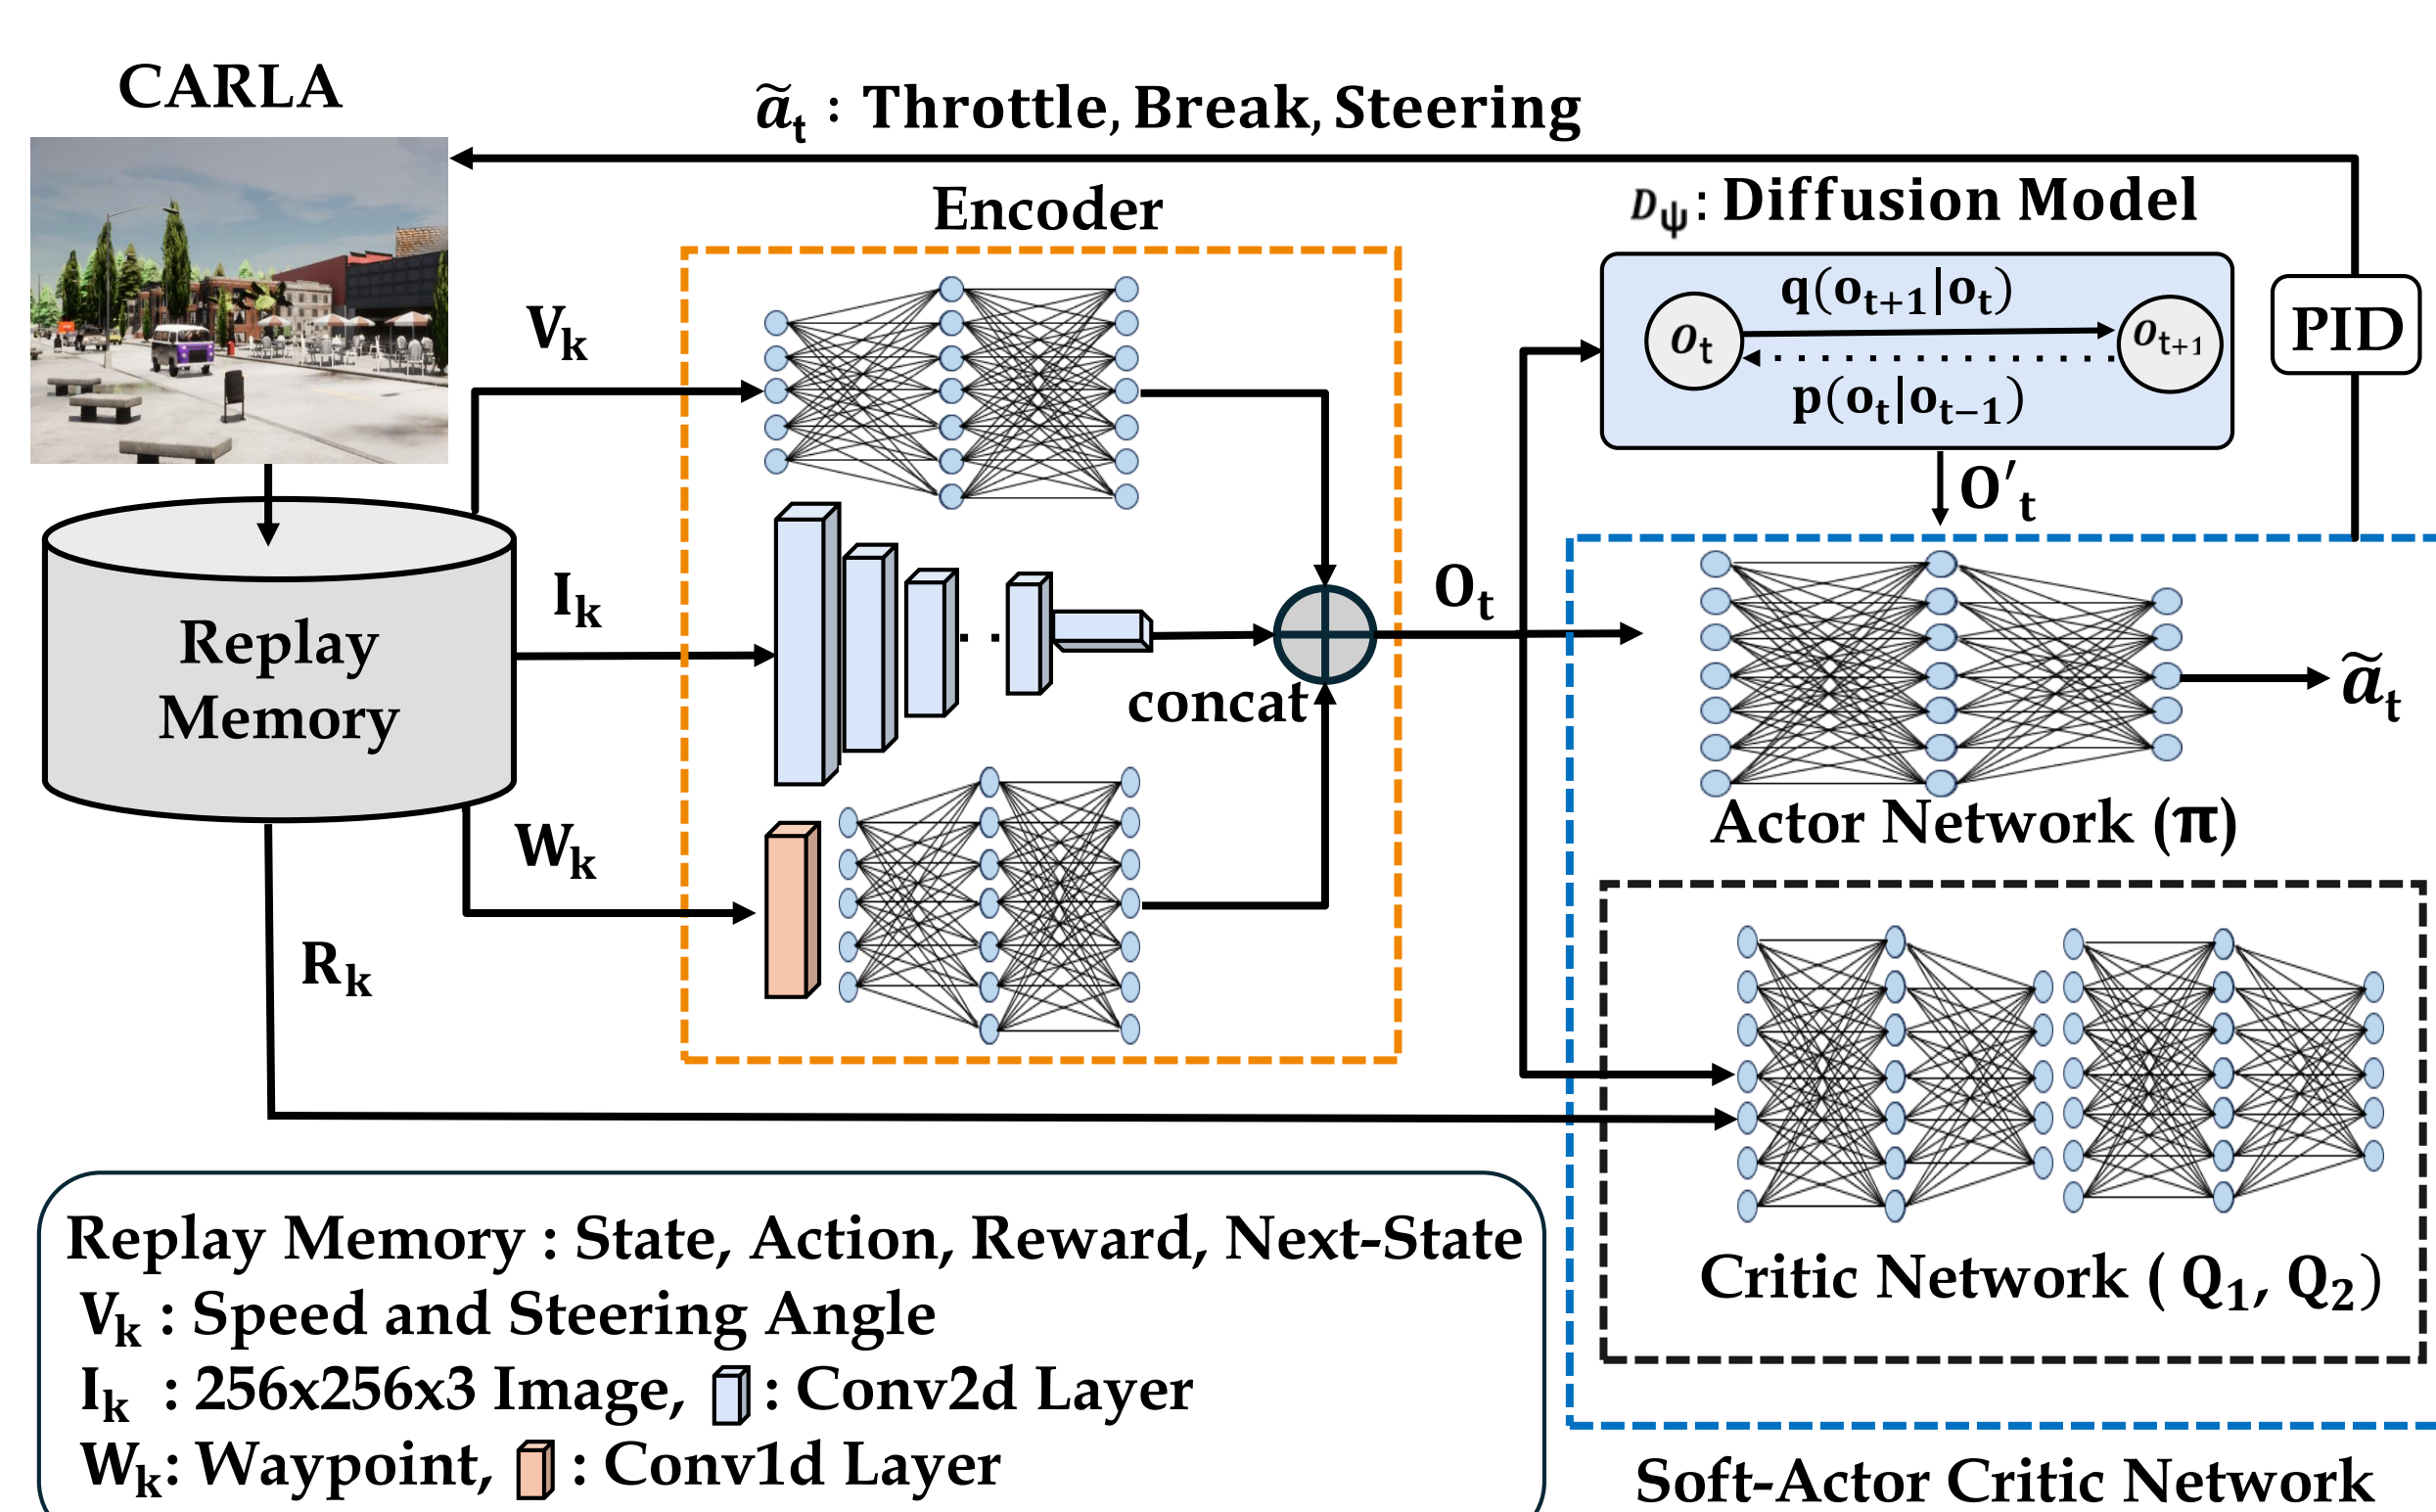
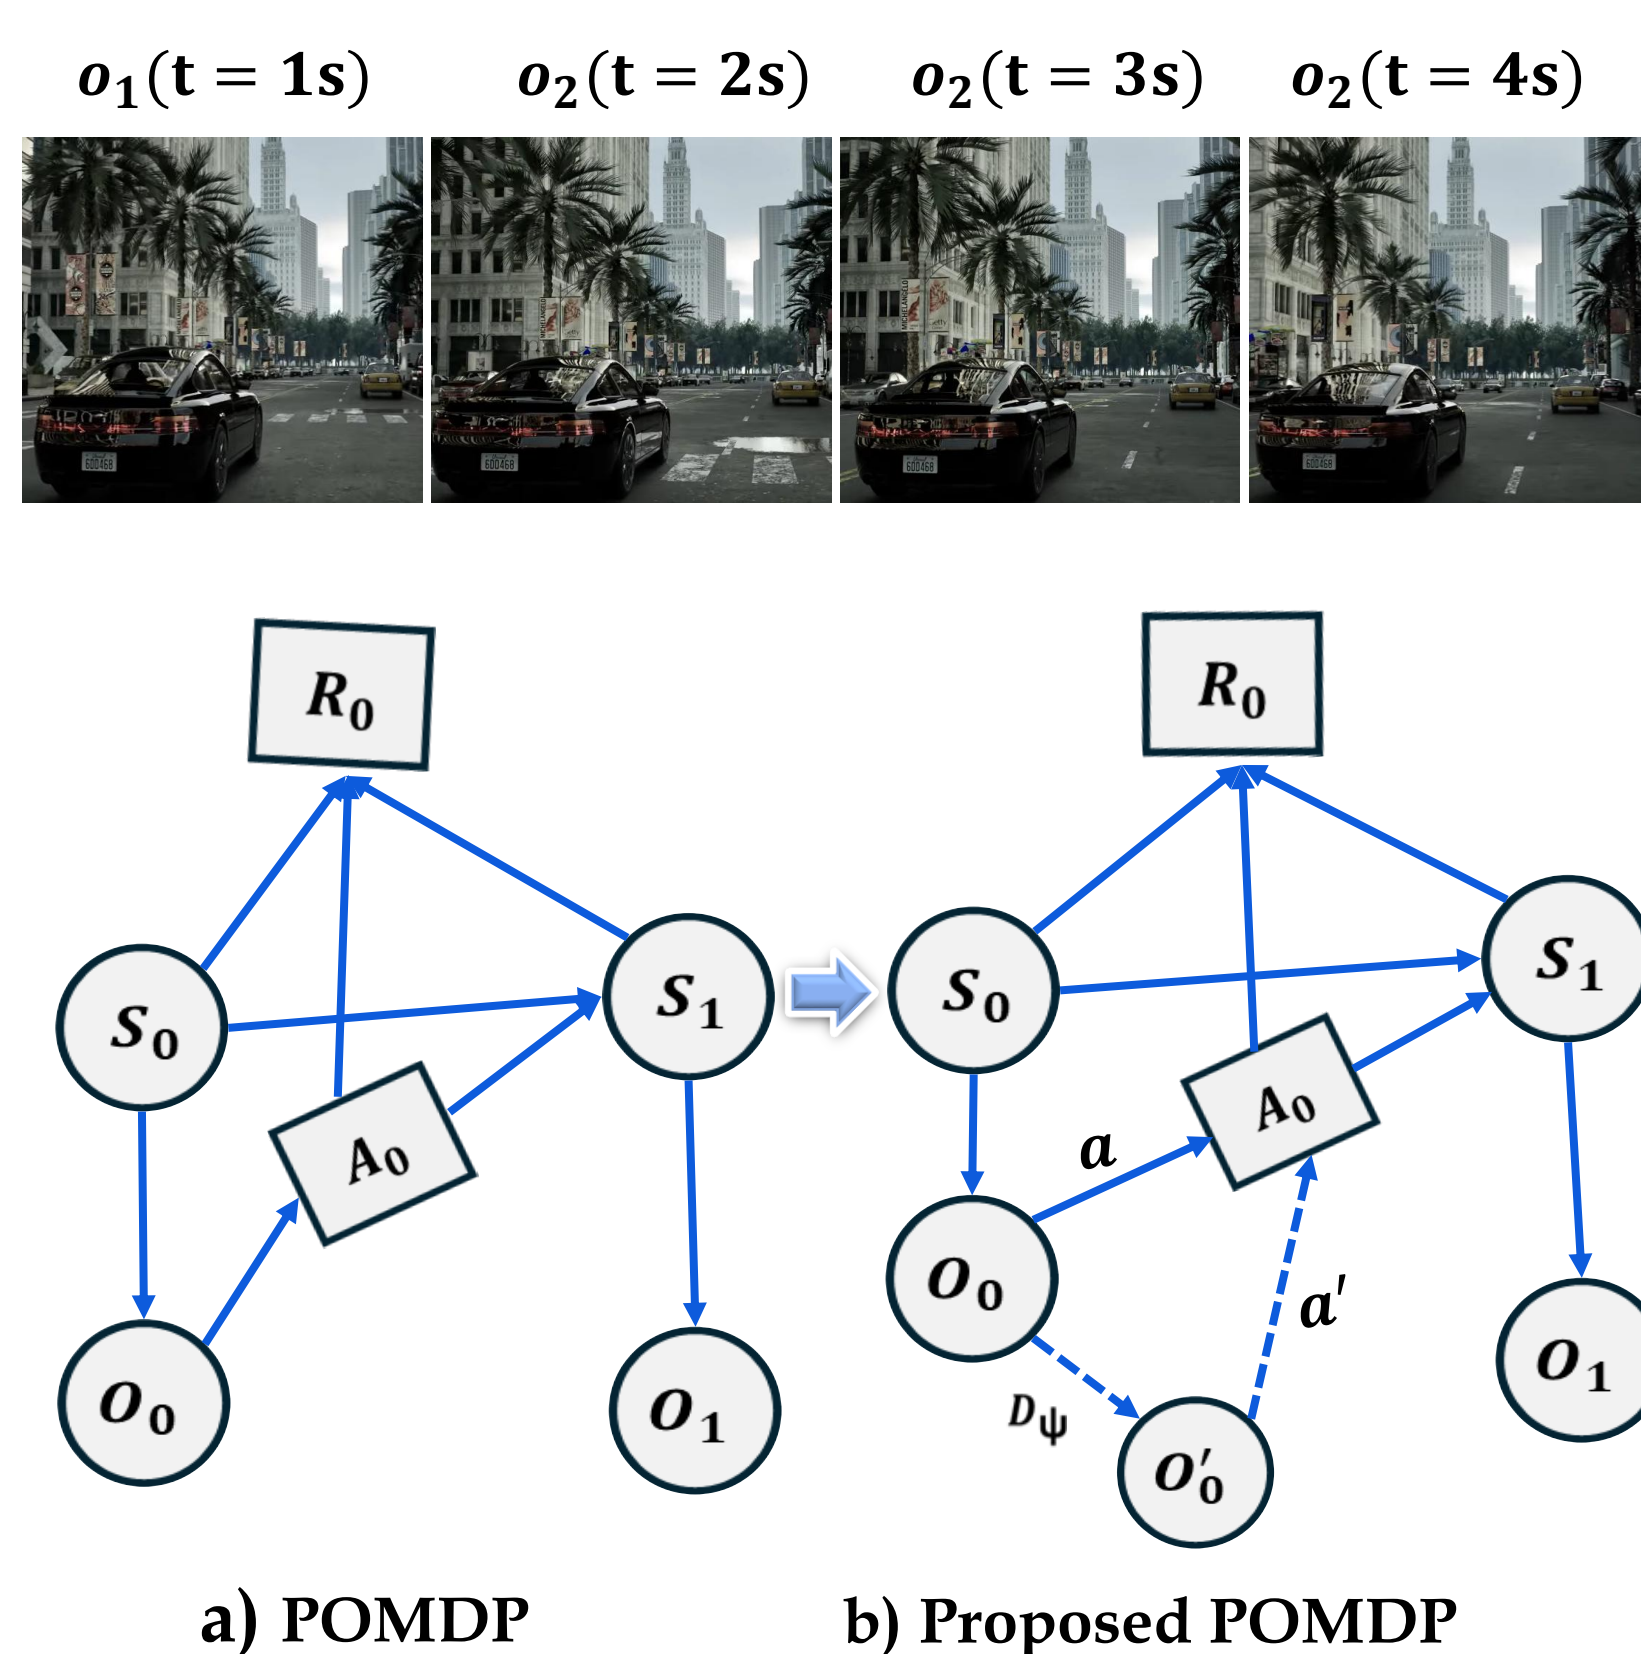


Figure 1: The DGDRL framework integrates encoder blocks, diffusion mode, and Soft Actor-Critic networks to address uncertainty and improve generalization in dynamic environments.

## Results




Figure 2: DGDRL performance comparison showing success rates (%) on CARLA NoCrash benchmark across two towns (T01, T02) and weather conditions (W01, W02): (a) T01, W01; (b) T01, W02; (c) T02, W01; (d) T02, W02




Figure 3: DGDRL training analysis across multiple metrics: (a-b) Latent observation comparison ( $o_t$  vs  $o'_t$ ) at 10K and 50K steps showing diffusion model's progressive enhancement in guiding policy network to handle environmental uncertainties, (c) Policy loss evolution over 500K training steps, and (d-e) Q-network loss curves comparing  $\lambda$  values (0.01 vs 0.1) during initial 50K steps to illustrate hyperparameter impact on learning dynamics.




Figure 4: DGDRL generalization analysis in Town02 dense setup. Rows 1-2: successful navigation through complex scenarios including underpasses. Bottom row: failure case where the model encounters a tunnel and halts, unable to make decisions. These examples highlight both model adaptability and limitations in specific environments.

## Conclusion and Future Work

- **Model:** The proposed Diffusion-Guided Deep Reinforcement Learning (DGDRL) model integrates diffusion models with Soft Actor-Critic under a modified POMDP framework.
- **Performance Benchmark:** 95.04% success rate across diverse environments (e.g., new town and weather)
- **Key Improvements:** Enhanced generalization to new task and improve pedestrian safety.
- **Future Work:** Targets complex scenarios such as tunnels, lane change, and broader domain applications.

## Acknowledgement

This work is partially supported by the DST-NSF Project and MeitY CPS Project, Govt. of India.

## References

- [1] Coelho et al., 2024] Daniel Coelho, Miguel Oliveira, and Vitor Santos. Rlfold: Reinforcement learning from online demonstrations in urban autonomous driving. In Proceedings of the AAAI Conference on Artificial Intelligence, volume 38, pages 11660–11668, 2024
- [2] Chekroun et al., 2023] Raphael Chekroun, Marin Toromanoff, Sascha Hornauer, and Fabien Moutarde. Gri: General reinforced imitation and its application to visionbased autonomous driving. Robotics, 12(5):127, 2023.
- [3] Zhao et al., 2022] Yimu Zhao, Kun Wu, Zhiyuan Xu, Zhengping Che, Qi Lu, Jian Tang, and Chi Harold Liu. Cadre: A cascade deep reinforcement learning framework for vision-based autonomous urban driving. In Proceedings of the AAAI, volume 36, pages 3481–3489, 2022.
- [4] Jain and Unhelkar, 2024] Abhinav Jain and Vaibhav Unhelkar. Go-dice: Goal-conditioned option-aware offline imitation learning via stationary distribution correction estimation. In Proceedings of the AAAI conference on artificial intelligence, volume 38, pages 12763–12772, 2024.
- [5] Pangarkar, 2025] Tajammul Pangarkar, Autonomous Vehicles Statistics 2025, <https://www.news.market.us/autonomous-vehicles-statistics>